

Online workshop on Generative AI

Hands-on session: Generative Adversarial Networks & Its Variants

Praveen K Chandaliya, Phd

Sardar Vallabhbhai National Institute of Technology, Surat
Department of Artificial Intelligence
pkc@aid.svnit.ac.in

March 9, 2025

Outline

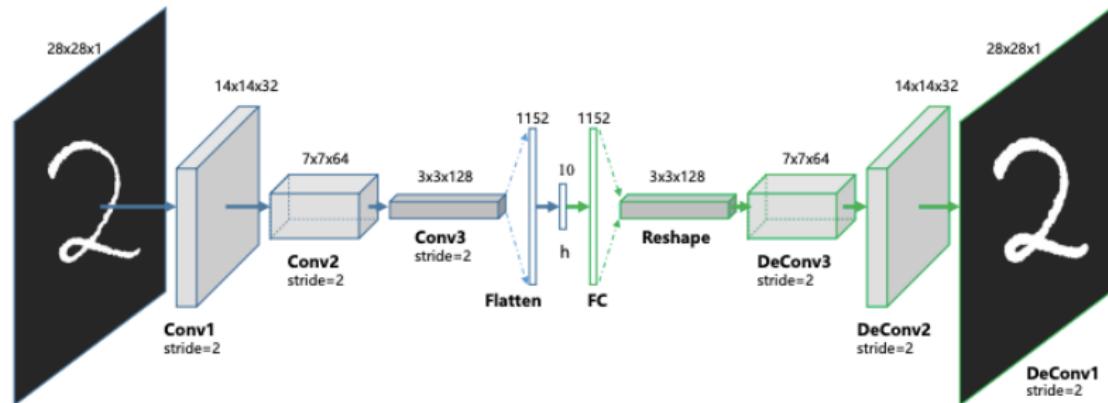
- 1 Prerequisites for GAN
- 2 Convolution AutoEncoder
- 3 Generative Adversarial Network
- 4 Objective function of GAN
- 5 Training GAN
- 6 GAN implementation on MNIST
- 7 Deep Convolution Generative Adversarial Network
- 8 GAN implementation on FaceDataset
- 9 Variants of GAN
- 10 Conditional Generative Adversarial Networks
- 11 LSGAN
- 12 Convolution AutoEncoder Face Dataset

Prerequisites for GAN

- Dataloader in PyTorch
- Convolution and Deconvloution Operation in Pytorch
- Sequential Convolution block
- MSE and Binary Cross entropy loss functions in Pytorch

STTP_Prerequisites.ipynb

Convolution Autoencoder

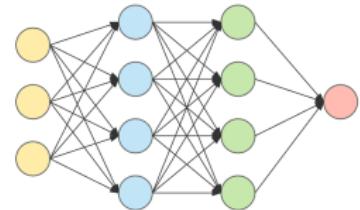


STTP_CNN_AE.ipynb

Generative Adversarial Network

Definition

- Generative
- Adversarial
- Networks



Generative Adversarial Network Architecture

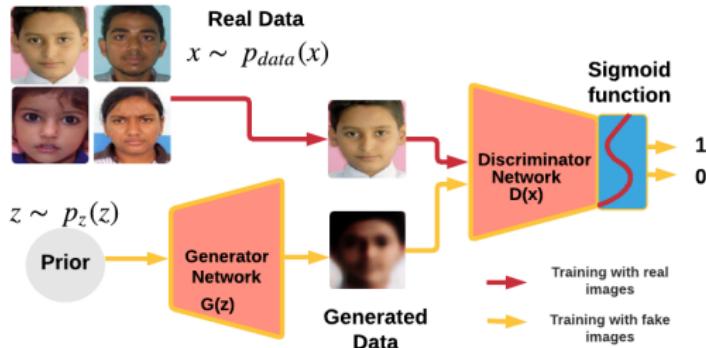


Figure 1: GAN Architecture

- Discriminator (D) distinguishing between real images and generated images and is responsible for classifying images as real(1) or fake (generated)(0).
- Generator (G) generating the images that fool the discriminator.

Loss Function

Binary Cross Entropy:

$$\mathcal{L}(\hat{y}, y) = y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}) \quad (1)$$

Discriminator on real (D) :

y = original, \hat{y} = generated image; The label for the data coming from $P_{data}(x)$ is $y = 1$ and $\hat{y} = D(x)$

$$\mathcal{L}(D(x), 1) = \log(D(x)) \quad (2)$$

Discriminator on fake (D):

Data coming from generator the label is $y = 0$ and $\hat{y} = D(G(z))$

$$\mathcal{L}(D(G(z)), 0) = \log(1 - D(G(z))) \quad (3)$$

Continue

Objective of the Discriminator is to correctly classify fake vs the real dataset. For this Equation 2 and 3 should be maximized.

$$\max_D \log(D(x)) + \log(1 - D(G(z))) \quad (4)$$

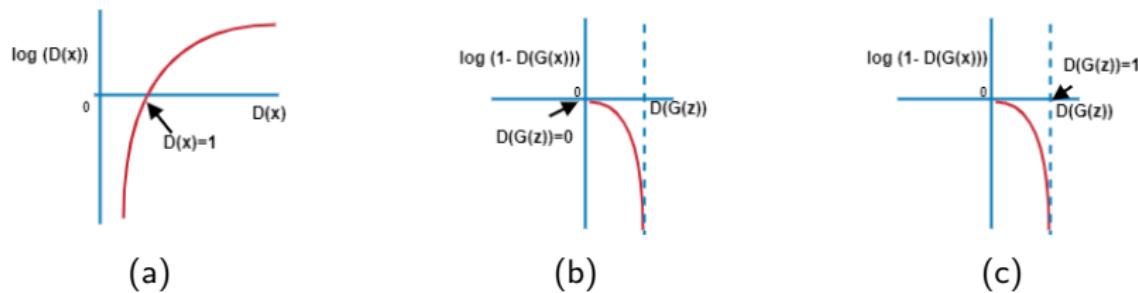


Figure 2: (a) and (b) Discriminator D Loss and (c) Generator G loss.

Objective of Generator to fool the discriminator. $D(G(z)) = 1$ then Generator fool the Discriminator by producing the probability as 1.

$$\min_G \log(1 - D(G(z))) \quad (5)$$

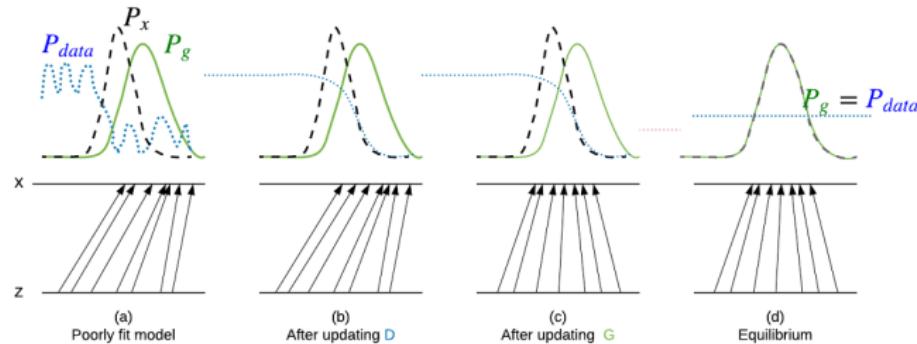
Mathematical Formulation

Training GAN is a two-player minimax game problem

- Discriminator D tries to maximize its classification accuracy.
- Generator G tries to minimize the discriminator classification accuracy.

$$\min_G \max_D V(D, G) = \log [D(x)] + \log [1 - D(G(z))] \quad (6)$$

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim P_{data}(x)} \log [D(x)] + \mathbb{E}_{z \sim P_z(z)} \log [1 - D(G(z))] \quad (7)$$



Training GAN

$$\min_G \max_D V(D, G)$$

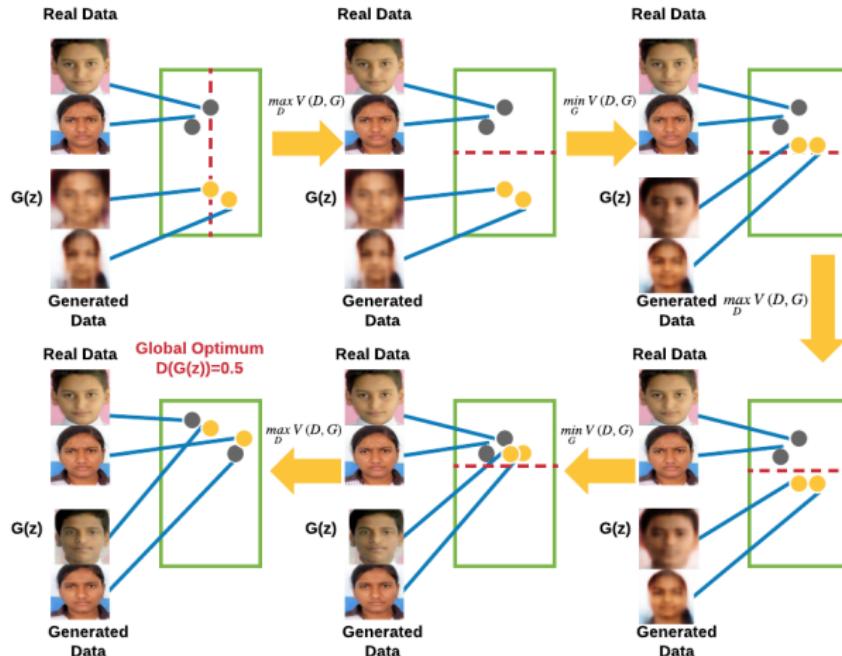


Figure 3: Training Steps in GAN

Algorithm

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

```
for number of training iterations do
  for  $k$  steps do
    • Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
    • Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
    • Update the discriminator by ascending its stochastic gradient:
```

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

```
end for
```

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

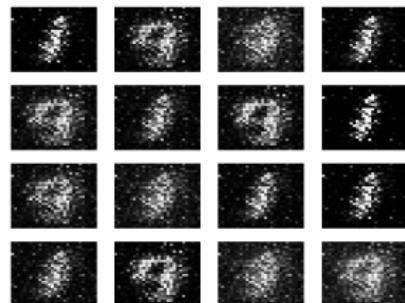
```
end for
```

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

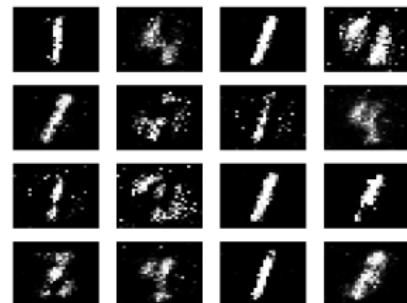
Figure 4: GAN Objective

Generative Adversarial Network implementation on colab (STTP_GAN.ipynb).

GAN Implementation Result: MNIST



(a) Epoch 5



(b) Epoch 20

Figure 5: GAN implementation result on MNIST

Challenges Faced by GANs

Training Instability

- ① Non-convergence: GANs often fail to converge due to the adversarial nature of their training, as the generator and discriminator compete with each other.
- ② Vanishing Gradients: When the discriminator becomes too good, the generator receives almost no gradient to improve itself.
- ③ Mode Collapse: The generator produces limited diversity, repeatedly generating similar outputs instead of a wide variety of data.
- ④ Sensitive Hyperparameters: Training requires careful tuning of hyperparameters like learning rates, batch sizes, and network architectures

Generator

- ① Hidden layers: Four 4x4 strided convolutional layers (1024, 512, 256, and 128 kernels, respectively) with ReLU.
- ② Output layer: 4x4 strided convolutional layer (4096 nodes = 64x64 size image) with Tanh.
- ③ Batch normalization is used except for output layer.

Discriminator

- ① Hidden layers: Four 4x4 convolutional layers (128, 256, 512, and 1024 kernels, respectively) with Leaky ReLU.
- ② Output layer: 4x4 convolutional layer (1 node) with Sigmoid.
- ③ Batch normalization is used except for 1st hidden layer and output layer.

Deep Convolution Generative Adversarial Network

Generator Discriminator

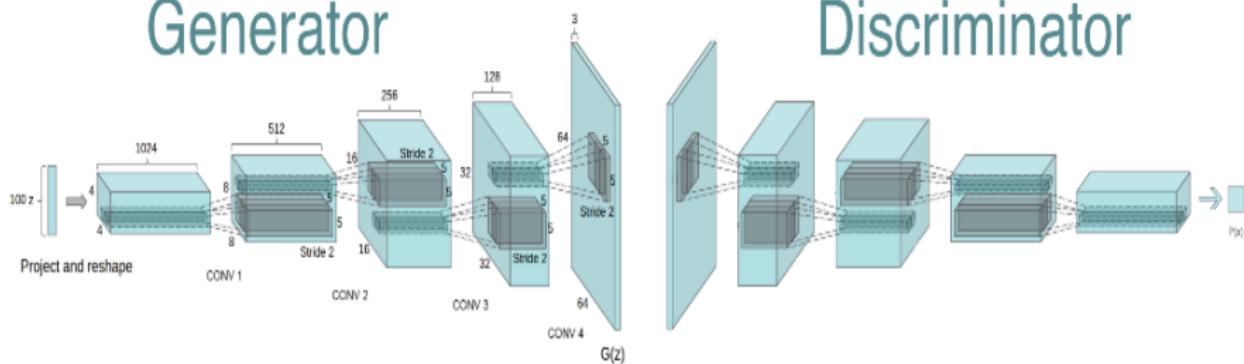
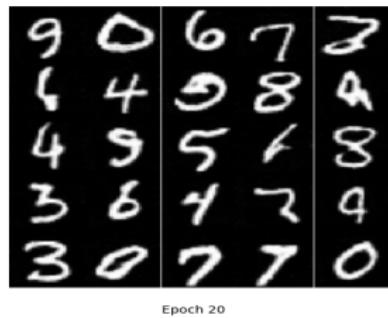


Figure 6: Deep Convolution Generative Adversarial Network ¹

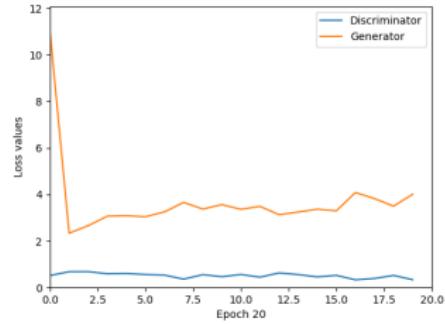
- 1 Replace any pooling layers with strided convolutions (discriminator) and fractional-strided convolutions (generator).
- 2 Batchnorm in both the generator and the discriminator.
- 3 ReLU activation in generator for all layers except for the output, which uses Tanh.
- 4 LeakyReLU activation in the discriminator for all layers.

¹ Soumith et al. "Unsupervised representation learning with deep convolutional generative adversarial networks."

DCGAN Implementation Result: MNIST



(a) MNIST Data



(b) Training Loss

Figure 7: DCGAN results on MNIST dataset (STTP_DCGAN_MNIST.ipynb)

DCGAN Implementation Result: Celeb



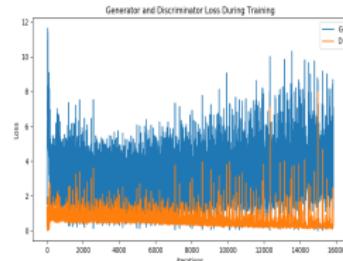
(a) Celeb Data



(b) epoch-1



(c) epoch-15



(d) Training Loss

Figure 8: DCGAN results on Celeb dataset

Variants of GAN structure

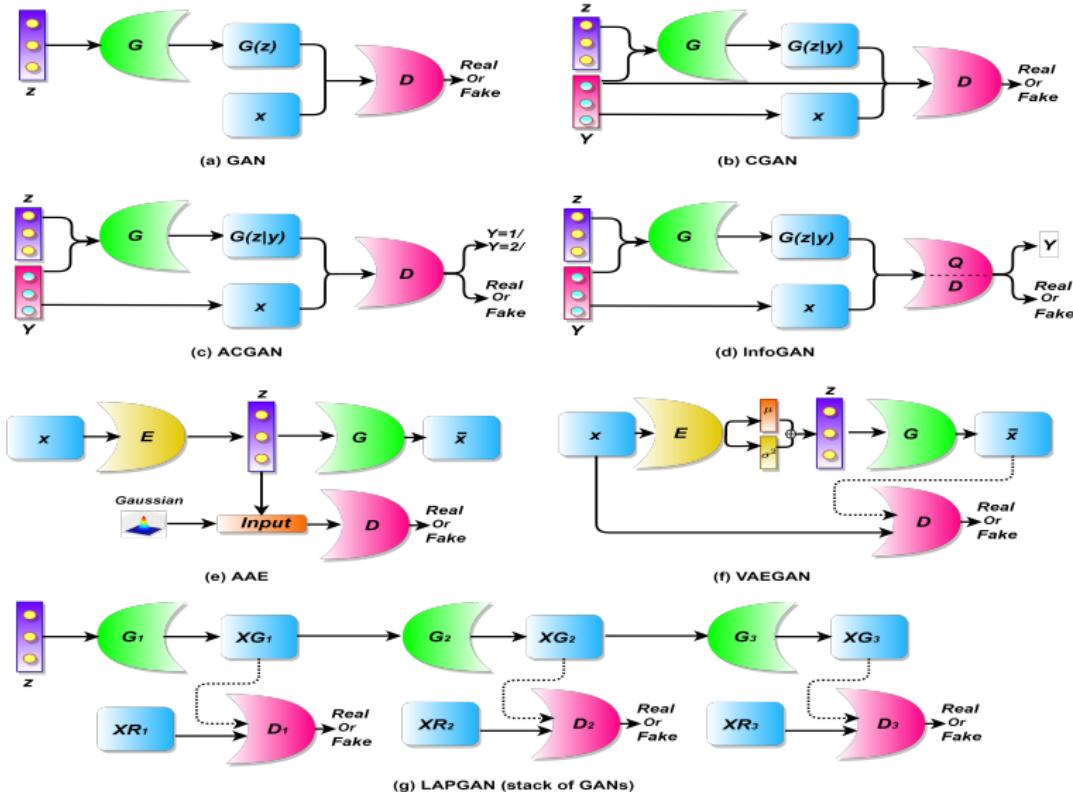


Figure 9: Variants of GAN structure

GAN Variants

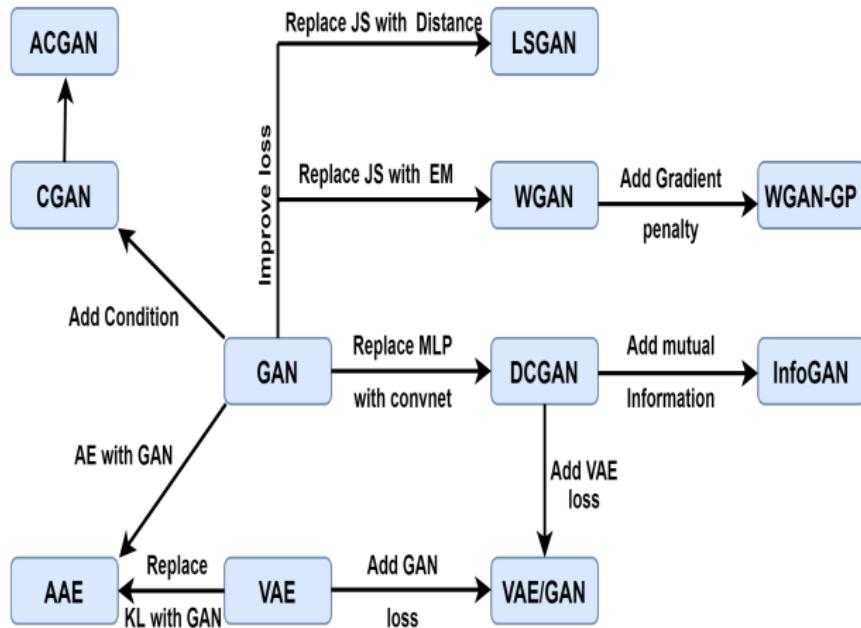


Figure 10: Variants of GAN structure

Conditional Generative Adversarial Network (cGAN)

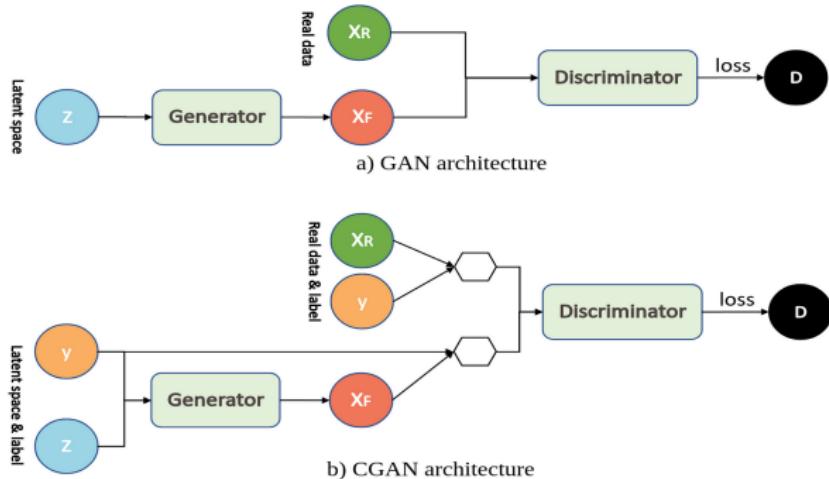


Figure 11: Conditional generative adversarial network

Conditional Generative Adversarial Network implementation on colab.

cGAN Result



Figure 12: cGAN

Least squares generative adversarial networks (LSGAN)

Least-square GAN

- For discriminator D has linear output

$$\min_D \frac{1}{2} E_{x \sim P_{data}} [(D(x) - b)^2] + \frac{1}{2} E_{x \sim P_G} [(D(x) - a)^2]$$

10

- For Generator

$$\min_D \frac{1}{2} E_{z \sim P_{data}} [(D(G(z)) - c)^2]$$

1

LSGAN:STTP_LSGAN.ipynb

Face Datasets

Table 1: Face datasets.

Dataset	Subjects	Images	Images Subject	Range	Avg.Age	Public
FGNET	82	1,002	6-18	0-69	16	Yes
Adience	2,958	26,580	NA	0-60	NA	Yes
Morph	13,000	55,134	2-53 (avg. 4.2)	16-77	42	Yes
CACD	2,000	163,446	22-139 (avg. 81.7)	16-62	31	Yes
IMDBWiki	20,284	523,051	avg. 25.79	0-100	38	Yes
UTKFace	NA	23,709	NA	0-116	33	Yes
AgeDB	568	16,488	avg. 29	1-101	50.3	Yes
AGFW-v2	36,325	27,688	-	10-64	-	Yes
ITWCC	754	7,990	3-37 (avg. 10.7)	0-32	13	Yes
CLF	919	3,682	2-6 (avg. 4.0)	2-18	8	No

Convolution Autoencoder with Face Dataset

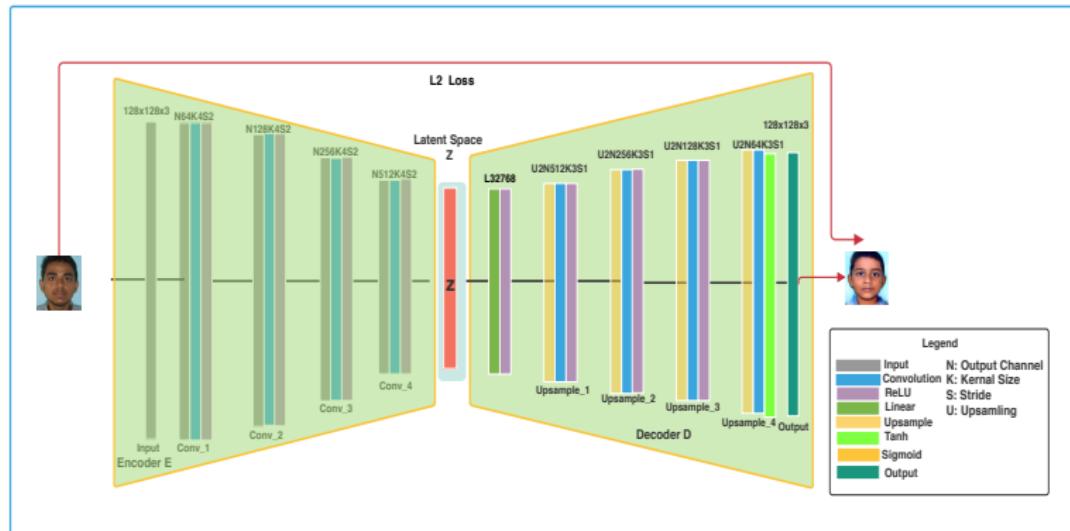


Figure 13: Convolution Autoencoder implementation on colab [STTP_CAE_Face.ipynb](#) ²

²Praveen Kumar Chandaliya, Neeta Nain, Conditional Perceptual Adversarial Variational Autoencoder, ICB2019, IEEE. ↗ ↘ ↙

Convolution Autoencoder Implementation Result



(a) Morph Data



(b) epoch-1



(c) epoch-15



(d) epoch-30

Figure 14: Convolution Autoencoder Implementation Results

- Praveen Kumar Chandaliya, Preyas Garg and Neeta Nain, " Retrieval of Facial Images Re-rendered with Natural Aging Effect using Child Facial Image and Age," (SITIS), Las Palmas de Gran Canaria, Spain, 2018, pp. 457-463
- Praveen Kumar Chandaliya, Neeta Nain "Conditional Perceptual Adversarial Variational Autoencoder for Age Progression and Regression on Children Face" ICB 2019. **B Level**,
- Praveen Kumar Chadnaliya, Vardhman, Mayank Harjani, Neeta Nain, "SCDAE: Ethnicity and Gender Alteration on CLF and UTKFace Dataset" CVIP 2019
- Praveen Kumar Chandaliya, Aditya Sinha, Neeta Nain "ChildFace: Gender Aware Child Face Aging" BIOSIG 2020 **B Level**.

References I



Neeta Nain Praveen Kumar Chandaliya.

Childgan: Face aging and rejuvenation to find missing children.
In *Pattern Recognition*, volume 129, page 108761, 2022.



P. Gallinari, Yann Lecun, S. Thiria, and F. Fogelman Soulie.

Memoires associatives distribuees: Une comparaison (distributed associative memories: A comparison).

In *Proceedings of COGNITIVA 87, Paris, La Villette, May 1987.*
Cesta-Afcet, 1987.



Diederik P. Kingma and Max Welling.

Auto-encoding variational bayes.

CoRR, abs/1312.6114, 2013.

References II

-  Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron C. Courville, and Yoshua Bengio. Generative adversarial nets.
In *NIPS*, 2014.
-  Mehdi Mirza and Simon Osindero. Conditional generative adversarial nets.
CoRR, abs/1411.1784, 2014.
-  Alireza Makhzani, Jonathon Shlens, Navdeep Jaitly, and Ian J. Goodfellow. Adversarial autoencoders.
CoRR, abs/1511.05644, 2015.

References III

 Alec Radford, Luke Metz, and Soumith Chintala.

Unsupervised representation learning with deep convolutional generative adversarial networks.

CoRR, abs/1511.06434, 2015.

 Debayan Deb, Neeta Nain, and Anil K. Jain.

Longitudinal study of child face recognition.

CoRR, abs/1711.03990, 2017.